

Turing Way-Aligned Readiness Snapshot

This is not an official certification. The Turing Way is public best-practice guidance, not a certification authority or institutional policy. This document is a transparent, evidence-based readiness snapshot produced by inspecting observable repository artifacts, scored against a ten-criterion rubric (0–10 each, 100 points total). It carries no legal, regulatory, or institutional weight, and it does not evaluate scientific correctness, security, or ethical/legal compliance beyond what is visible in the repository.

- **Repository:** `stjude-biohackathon/KIDS25-Team18`
- **Reviewed commit:** `75406c545ec7a6c6cc6e390e2393834828deee08` (branch `main`, dated 2025-10-07)
- **Review date:** 2026-09-18
- **Scope:** Full public repository contents at the reviewed commit — source code (R, Python, Nextflow), documentation, licensing, and committed data artifacts. No private CI logs, issue/discussion history, or non-public metadata were available to inspect.
- **Confidence:** Moderate–High. Findings are based on direct inspection of the full working tree (all top-level files/directories), `git log / git tag` output, and file contents (README, LICENSE, `methynet/README`, `nextflow.config`, `main_nextflow.nf`). Confidence is not "High" because GitHub-hosted metadata not present in the local checkout (e.g., branch protection rules, Issues/Discussions activity, GitHub Actions run history beyond the absence of workflow files) could not be directly inspected.

Score summary

#	Criterion	Score (0–10)
1	Project purpose and scope	0
2	Version control and provenance	4
3	Open collaboration	0
4	Reproducible environments	2
5	Data and workflow provenance	2
6	Testing and validation	0
7	Automation and continuous integration	0
8	Documentation and usability	2
9	Licensing, attribution, and responsible reuse	6

#	Criterion	Score (0–10)
10	Ethics, accessibility, and sustainability	0
Total		16 / 100

Criterion detail

1. Project purpose and scope — 0/10

Evidence: `README.md` contains only the single line `# KIDS25-Team18` (15 bytes total). No description of the research question, purpose, scope, intended users, or known limitations exists anywhere in the repository. **Check:** `cat README.md → # KIDS25-Team18`. **Recommendation:** Add a README section stating the project's purpose (e.g. the DepMap/methylation prediction goal implied by the code), intended audience, and known limitations. **Turing Way guidance:** [Subprojects](#)

2. Version control and provenance — 4/10

Evidence: `git log` shows 46 commits on `main` with a GitHub remote (`stjude-biohackathon/KIDS25-Team18`). Messages are mostly descriptive (e.g. "Add test_genes", "KNN performance R plot") but some are informal ("post event", "accept ty's changes"). `git tag` returns no tags — there are no versioned releases. **Check:** `git log --oneline | wc -l → 46; git tag → (empty)`. **Recommendation:** Adopt tagged releases (e.g. semantic version tags) at milestones, and use more consistently descriptive commit messages. **Turing Way guidance:** [Contributors Record](#)

3. Open collaboration — 0/10

Evidence: No `CONTRIBUTING.md`, `CODE_OF_CONDUCT.md`, issue/PR templates, or `.github/` directory of any kind exist in the repository. No maintainership information is documented. **Check:** `find . -iname "CONTRIBUTING*" -o -iname "CODE_OF_CONDUCT*" → no results; .github/ directory absent`. **Recommendation:** Add a `CONTRIBUTING.md` describing how to propose changes, a `CODE_OF_CONDUCT.md`, and issue templates to set collaboration expectations. **Turing Way guidance:** [Collaboration](#)

4. Reproducible environments — 2/10

Evidence: `methynet/README` documents manual `conda/pip` install steps for one subcomponent (`conda create -n methynet python=3.6, pip install methynet`), but no `requirements.txt`, `environment.yml`, `renv.lock`, or Dockerfile exists anywhere in the repository for the R scripts (`build_model*.R`), the Nextflow pipeline, or the `tabpfn_investigation/bin` Python scripts. **Check:** `find . -iname "*requirements*" -o -iname "*environment*.yml" -o -iname "*renv*" → no results outside the one narrative README`.

Recommendation: Add pinned dependency manifests (e.g. `renv.lock` for R, `environment.yml/requirements.txt` for Python) and, ideally, a container definition for the Nextflow pipeline. **Turing Way guidance:** [Data Usage Statement](#)

5. Data and workflow provenance — 2/10

Evidence: The repository root commits large data/result artifacts directly (`depmap_predictions.csv` ~7.6 MB, `importance_df3_MYCN_XGB.csv`, and four timestamped `results_*/` output directories) with no accompanying documentation of their schema, generation date/parameters, or provenance. A Nextflow workflow (`main_nextflow.nf`, `nextflow.config`, `nextflow_modules/pfn_quick.nf`) does describe pipeline steps at a basic level. **Check:** `ls results_*` and `wc -l depmap_predictions.csv` (17,917 lines) with no adjacent data dictionary.

Recommendation: Document each data file's origin, generation command, and schema (e.g. a `DATA.md`), and consider excluding large generated artifacts from version control in favor of a data-management plan. **Turing Way guidance:** [Data Usage Statement](#)

6. Testing and validation — 0/10

Evidence: No test framework, test directory, or assertions exist. Files named `test_genes.csv`, `small_test_input.txt`, and `large_test_jobs.txt` are ad-hoc input data, not automated tests, and there is no documented procedure for validating results. **Check:** repository-wide search for test frameworks (pytest, testthat, etc.) found none. **Recommendation:** Add automated tests (e.g. `testthat` for R scripts, `pytest` for Python) with representative fixtures and document how to run them. **Turing Way guidance:** [Glossary](#)

7. Automation and continuous integration — 0/10

Evidence: No `.github/workflows/` directory or any other CI configuration (e.g. `.gitlab-ci.yml`, `Jenkinsfile`) exists in the repository. **Check:** `find .github/workflows` → not found. **Recommendation:** Add CI (e.g. GitHub Actions) to run linting/tests on pull requests and surface failures to contributors. **Turing Way guidance:** [Legal Disclaimer](#)

8. Documentation and usability — 2/10

Evidence: The top-level `README.md` is a one-line title with no usage instructions. Only `methynet/README` documents setup/run instructions, and only for that one subcomponent; the R model-building scripts, Nextflow pipeline, and `tabpfn_investigation` scripts have no onboarding documentation. **Check:** `wc -c README.md` → 15 bytes; only one of several subcomponents has any README.

Recommendation: Expand the top-level README with an overview, setup instructions for each subcomponent, and usage examples. **Turing Way guidance:** [Afterword](#)

9. Licensing, attribution, and responsible reuse — 6/10

Evidence: A clear MIT LICENSE file is present at the repository root, attributed to "St. Jude Children's Research Hospital BioHackathon." No CITATION.cff or explicit third-party attribution/reuse guidance exists for bundled third-party tools referenced in the code (e.g. methylnet, pymethylprocess, tabpfn). **Check:** cat LICENSE confirms MIT terms; no CITATION.cff found. **Recommendation:** Add a CITATION.cff for citing the repository and note third-party tool licenses/attribution explicitly. **Turing Way guidance:** [Legal Disclaimer](#)

10. Ethics, accessibility, and sustainability — 0/10

Evidence: No files address data ethics/privacy (relevant given genomic and cell-line data usage), accessibility of project outputs, or a sustainability/maintenance/archival plan.

Check: Repository-wide search for ethics/privacy/accessibility/ maintenance documentation → none found. **Recommendation:** Add a short statement on data governance/ethics for the genomic data used, and a maintenance/archival plan (even informal, e.g. who maintains this post-hackathon). **Turing Way guidance:** [Data Usage Statement](#)

Limitations of this snapshot

- This is a readiness snapshot based on static repository inspection at one commit, not an official Turing Way certification, security review, or scientific validation.
- GitHub-side collaboration signals (Issues, Discussions, PR review history, branch protection) were not accessible from the local checkout and were not scored as present or absent beyond what committed files show.
- Scores reflect only observable, verifiable evidence in the repository at the time of review; they may change as the repository evolves.

Total score: 16 / 100